

Decision Trees

K. Gibert⁽¹⁾

*(¹)Department of Statistics and Operation Research
Knowledge Engineering and Machine Learning group at
Intelligent Data Science and Artificial Intelligence Research Center
Universitat Politècnica de Catalunya, Barcelona*

karina.gibert@upc.edu

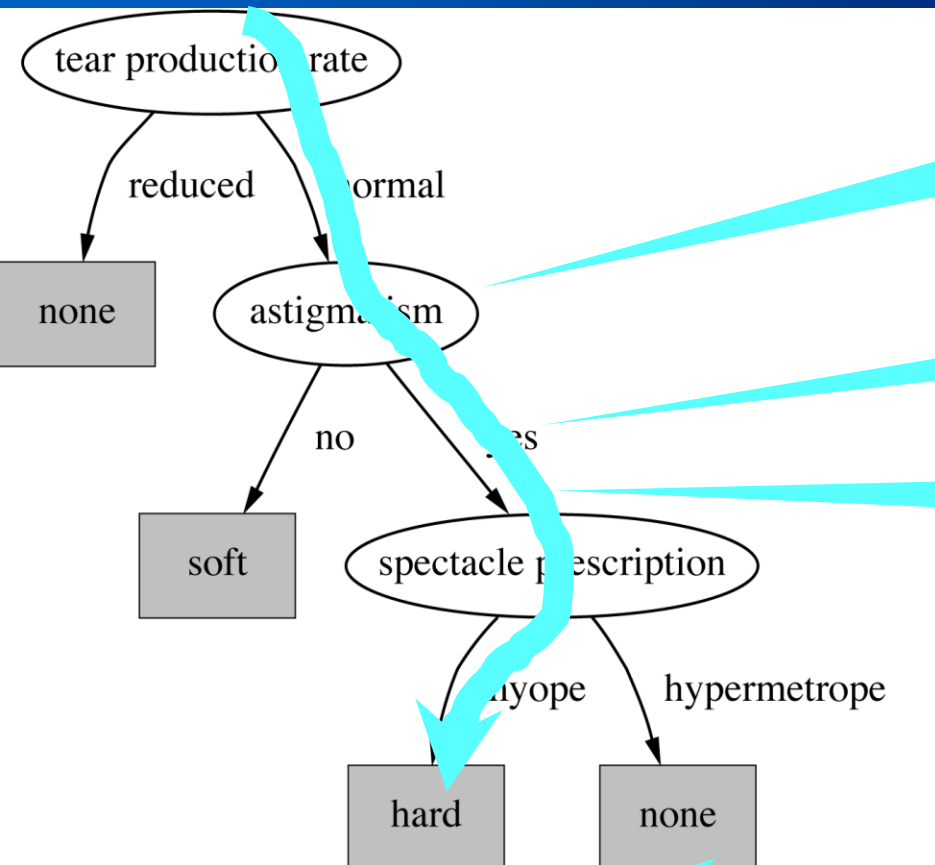
<https://www.eio.upc.edu/homepages/karina>

Decision Trees

- **Predict a class variable**
- Visual output self-understandable

Decision trees

Model the process of deciding to which class belongs a new example of the domain



Internal nodes: Variables

Branches: modalities

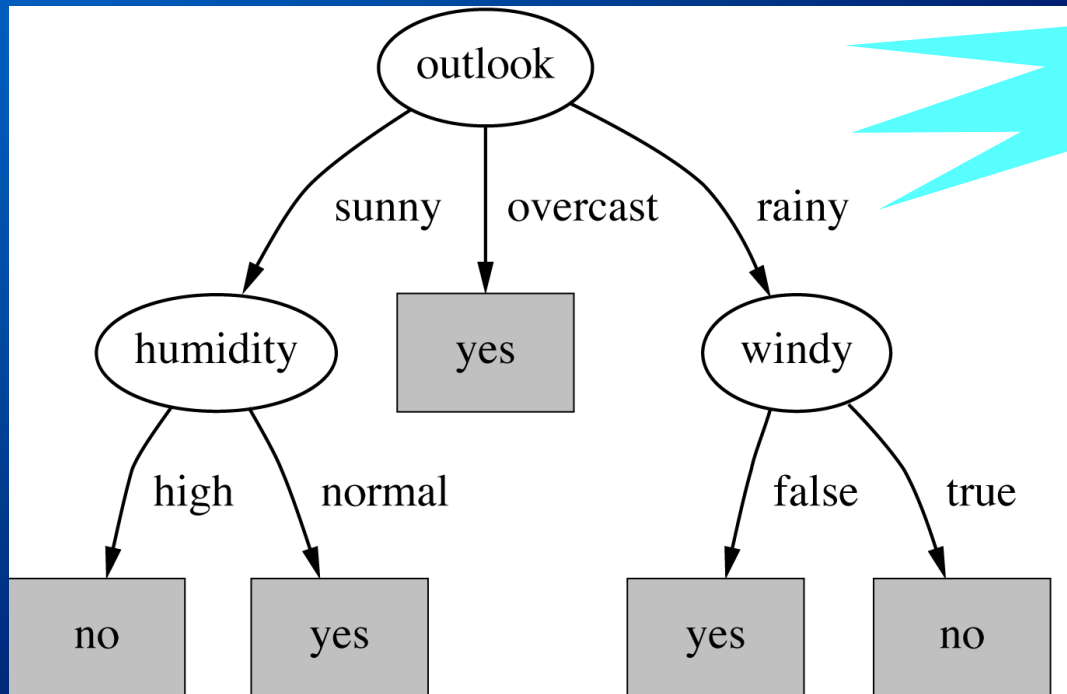
Paths: Inducted decisions

Associated classification rules

If tear prod rate=normal **and**
astigmatism= yes **and**
spectacle prescription=myope
then hard lent

Leafs: Predicted class

Decision Tree: example 2



Easy to understand

Construction:

- Top-down strategy
- Which splitting variable?
- Which branches?
- How to label leafs?

Decision tree for the weather data

Leafs:
Voting

Decision Trees

□ Many models:

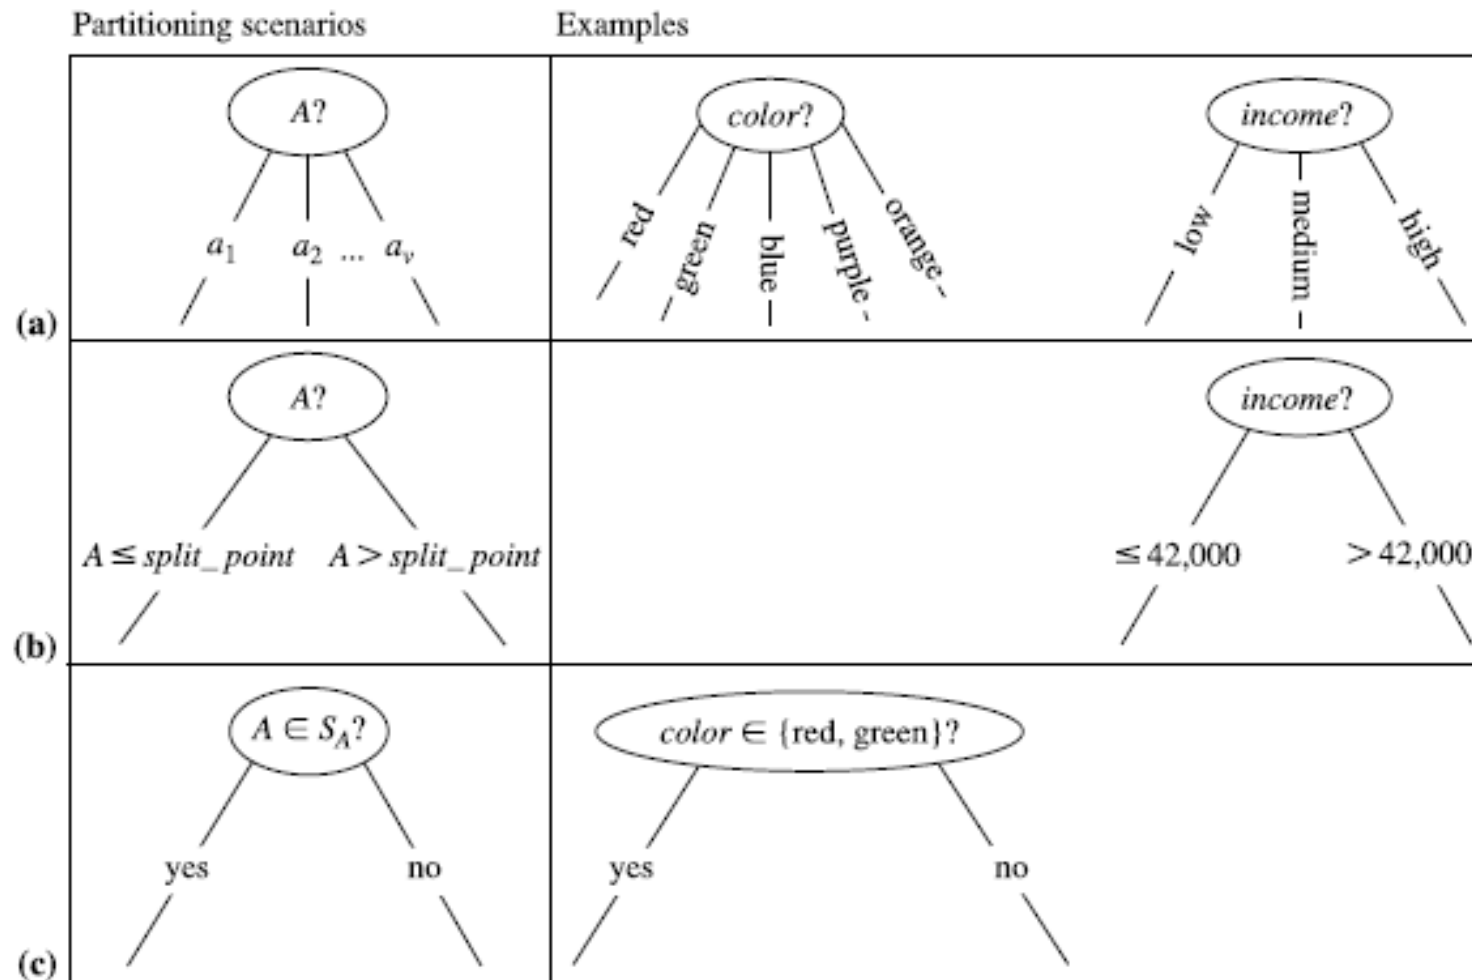
- For numerical/qualitative response variables
- For binary, nominal, ordinal, numerical explanatory variables
- Binary tree/multiway-tree
- Split criterion (information gain, gini index...)
- Stop criterion (pre-pruning, post-pruning)

ID3, CART (previous AID, CHAID...)

Decision Trees

- General algorithm:
 - Set all individuals at the root node
 - Optimize split in children nodes
 - For every child: decide between
 - Continue splitting (repeating process)
 - Stop

Decision Trees



ID3: Entropy & Information Gain

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i), \quad (8.1)$$

where p_i is the nonzero probability that an arbitrary tuple in D belongs to class C_i and is estimated by $|C_{i,D}|/|D|$. A log function to the base 2 is used, because the information is encoded in bits. $Info(D)$ is just the average amount of information needed to identify the class label of a tuple in D . Note that, at this point, the information we have is based solely on the proportions of tuples of each class. $Info(D)$ is also known as the **entropy** of D .

How much more information would we still need (after the partitioning) to arrive at an exact classification? This amount is measured by

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j). \quad (8.2)$$

The term $\frac{|D_j|}{|D|}$ acts as the weight of the j th partition. $Info_A(D)$ is the expected information required to classify a tuple from D based on the partitioning by A . The smaller the expected information (still) required, the greater the purity of the partitions.

ID3: Entropy & Information Gain

Information gain is defined as the difference between the original information requirement (i.e., based on just the proportion of classes) and the new requirement (i.e., obtained after partitioning on A). That is,

$$Gain(A) = Info(D) - Info_A(D). \quad (8.3)$$

In other words, $Gain(A)$ tells us how much would be gained by branching on A . It is the expected reduction in the information requirement caused by knowing the value of A . The attribute A with the highest information gain, $Gain(A)$, is chosen as the splitting attribute at node N . This is equivalent to saying that we want to partition on the attribute A that would do the “best classification,” so that the amount of information still required to finish classifying the tuples is minimal (i.e., minimum $Info_A(D)$).

ID3: Entropy & Information Gain

Table 8.1 Class-Labeled Training Tuples from the *AllElectronics* Customer Database

<i>RID</i>	<i>age</i>	<i>income</i>	<i>student</i>	<i>credit_rating</i>	<i>Class: buys_computer</i>
1	youth	high	no	fair	no
2	youth	high	no	excellent	no
3	middle_aged	high	no	fair	yes
4	senior	medium	no	fair	yes
5	senior	low	yes	fair	yes
6	senior	low	yes	excellent	no
7	middle_aged	low	yes	excellent	yes
8	youth	medium	no	fair	no
9	youth	low	yes	fair	yes
10	senior	medium	yes	fair	yes
11	youth	medium	yes	excellent	yes
12	middle_aged	medium	no	excellent	yes
13	middle_aged	high	yes	fair	yes
14	senior	medium	no	excellent	no

ID3: Entropy & Information Gain

$$Info(D) = -\frac{9}{14} \log_2 \left(\frac{9}{14} \right) - \frac{5}{14} \log_2 \left(\frac{5}{14} \right) = 0.940 \text{ bits.}$$

$$Gain(age) = Info(D) - Info_{age}(D) = 0.940 - 0.694 = 0.246 \text{ bits.}$$

$$Gain(income) = 0.029 \text{ bits, } Gain(student) = 0.151 \text{ bits,}$$

$$Gain(credit_rating) = 0.048 \text{ bits.}$$

ID3 algorithm



- **ID3** \equiv Induction Decision Tree [Quinlan, 1979]
- Machine Learning Technique
- Decision Tree Induction
- Top-Down strategy
- Qualitative variables
- Multi-way trees
- Don't repeat variables
- From a set of **examples (instances + belonging class)** build up the *best* decision tree which explains the instances

Criteria

Compactness → the more compact the better

Goodness → predictive/discriminant ability

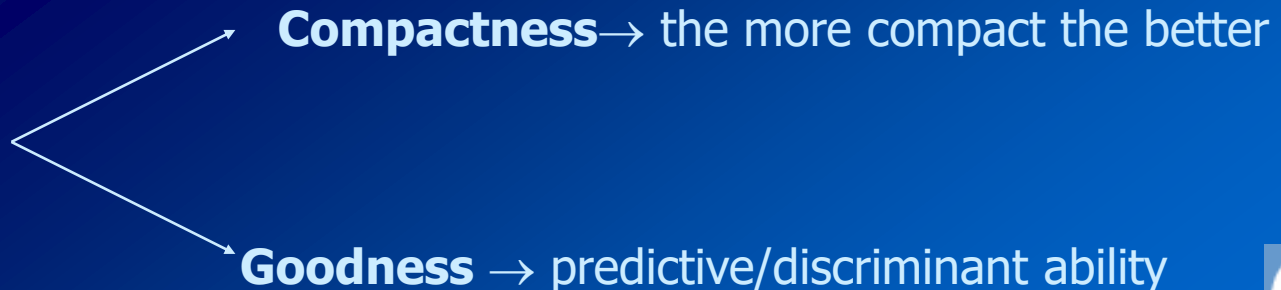
© K. Gibert

C4.5 algorithm



- C4.5: Refines ID3 [Quinlan, 1986]
- Machine Learning Technique
- Decision Tree Induction
- Top-Down strategy
- Qualitative variables
- Multi-way trees
- Uses impurity ratio
- Heuristics for pruning
- From a set of **examples (instances + belonging class)** build up the *best* decision tree which explains the instances

Criteria



C4.5: Gain Ratio

Gain Ratio

The information gain measure is biased toward tests with many outcomes. That is, it prefers to select attributes having a large number of values. For example, consider an attribute that acts as a unique identifier such as *product_ID*. A split on *product_ID* would result in a large number of partitions (as many as there are values), each one containing just one tuple. Because each partition is pure, the information required to classify data set D based on this partitioning would be $Info_{product_ID}(D) = 0$. Therefore, the information gained by partitioning on this attribute is maximal. Clearly, such a partitioning is useless for classification.

C4.5, a successor of ID3, uses an extension to information gain known as *gain ratio*, which attempts to overcome this bias. It applies a kind of normalization to information gain using a “split information” value defined analogously with $Info(D)$ as

$$SplitInfo_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \times \log_2 \left(\frac{|D_j|}{|D|} \right). \quad (8.5)$$

C4.5: Gain Ratio

$$\text{GainRatio}(A) = \frac{\text{Gain}(A)}{\text{SplitInfo}_A(D)}. \quad (8.6)$$

The attribute with the maximum gain ratio is selected as the splitting attribute. Note, however, that as the split information approaches 0, the ratio becomes unstable. A constraint is added to avoid this, whereby the information gain of the test selected must be large—at least as great as the average gain over all tests examined.

ID3

Which variable in decision node?

It is a *greedy* (*voraz*) algorithm

the *best* attribute at each step

is the most discriminant one (potentially more useful)

Maximize a certain $G(I, X)$

ID3: Notation

- $I \equiv$ set of examples $\equiv i = \{i_1.. i_n\}$
 - $X \equiv$ variable with values $\{x_1.. x_{nx}\}$
 - $P \equiv$ set of classes $\equiv \{c_1..... c_p\}$
 - $\# \equiv$ Cardinality
- $X^{-1}(x_i) \equiv$ set of elements in I with $X=x_i$

ID3: splitting criteria

- Select variable X which *maximizes* the information gain

$$G(I, X) = H(I) - H(I, X)$$

where

P response variable (class)

Divide by $H(I)$
Avoid effect of n_x

$$H(I) = - \sum_{c \in P} p(c|I) * \log_2 p(c|I)$$

Information
of parent node

$$p(c|I) = \#c / \#I$$

Empirical probability of
being in class C

Entropy
of splitting
by X

$$H(I, X) = \sum_{i=1:n_x} p_X(x_i) * H(X^{-1}(x_i))$$

Probability that one
example has the
value x_i for the
variable X

$$p_X(x_i) = \#X^{-1}(x_i) / \#I$$

Termination criteria

- All instances in a node belong to same class
- A maximum depth of the tree is achieved
(maxdepth control parameter in R)
- Splits do not open a minimum number of branches
(only for n-ary trees)
- Next split produces too small leafs
(minsplit control parameter in R)
- Improvement of termination function is not enough
(cp control parameter in R, termination function = splitting criterion)

Post-prunning criteria

- Avoids overfitting and redundancy
- Evaluate the tree with/without the branch
 - If performance is maintained, prune
 - Performance can be measured in different ways
 - Different error functions
 - Statistical tests

Algorithm BuildDecisionTree(\mathcal{D} , \mathcal{A})

input: training data \mathcal{D} , set \mathcal{A} of available attributes

output: a decision tree matching \mathcal{D} , using all or a subset of \mathcal{A}

```
1  if all elements in  $\mathcal{D}$  belong to one class
2      return node with corresponding class label
3  elseif  $\mathcal{A} = \emptyset$ 
4      return node with majority class label in  $\mathcal{D}$ 
5  else
6      select attribute  $A \in \mathcal{A}$  which best classifies  $\mathcal{D}$ 
7      create new node holding decision attribute  $A$ 
8      for each split  $v_A$  of  $A$ 
9          add new branch below with corresponding test for this split
10         create  $\mathcal{D}(v_A) \subset \mathcal{D}$  for which split condition holds
11         if  $\mathcal{D}(v_A) = \emptyset$ 
12             return node with majority class label in  $\mathcal{D}$ 
13         else
14             add subtree returned by calling
15                 BuildDecisionTree( $\mathcal{D}(v_A)$ ,  $\mathcal{A} \setminus \{A\}$ )
16         endif
17     endfor
18     return node.
19 endif
```

ID3: EXAMPLE (9)

	Eye Colour	Hair Colour	Height	Class
E1	Blue	Blonde	Tall	C+
E2	Blue	Brown	Medium	C+
E3	Brown	Brown	Medium	C-
E4	Green	Brown	Medium	C-
E5	Green	Brown	Tall	C+
E6	Brown	Brown	Low	C-
E7	Green	Blonde	Low	C-
E8	Blue	Brown	Medium	C+

ID3: example (10)

$$I = \{E1, \dots, E8\}$$

$$C+ = \{E1, E2, E5, E8\}$$

$$C- = \{E3, E4, E6, E7\}$$

$$G(I, \text{Eye-Color}) = H(I) - H(I, \text{Eye-Color})$$

$$H(I) = -\sum_{C_i \in C} p(C_i | I) * \log_2 p(C_i | I) = -p(C+ | I) * \log_2 p(C+ | I) - p(C- | I) * \log_2 p(C- | I)$$

$$p(C+ | I) = p(C- | I) = 1/2$$

$$H(I) = -1/2 \log_2 1/2 - 1/2 \log_2 1/2 = 1$$

$$H(I, \text{Eye-Color}) = \sum_{x_i} p_X(x_i) * H(X^{-1}(x_i)) = \sum_{x_i} p_{\text{Eye-Color}}(x_i) * H(\text{Eye-Color}^{-1}(x_i))$$

$$P_X(x_1) = P_{\text{Eye-Color}}(\text{Blue}) = 3/8$$

$$\begin{aligned} H(\text{Eye-Colour}^{-1}(\text{Blue})) &= -p(C+ | \text{Blue}) * \log_2 p(C+ | \text{Blue}) - p(C- | \text{Blue}) * \log_2 p(C- | \text{Blue}) \\ &= -1 \log_2 1 - 0 \log_2 0 \end{aligned}$$

$$\begin{aligned} H(I, \text{Eye-colour}) &= \frac{3}{8} (-1 \log_2 1 - 0 \log_2 0) + \frac{2}{8} (-0 \log_2 0 - 1 \log_2 1) + \\ &\quad \frac{3}{8} (-1/3 \log_2 1/3 - 2/3 \log_2 2/3) = 0.344 \end{aligned}$$

ID3: example (10)

$$H(I) = -1/2 \log_2 1/2 - 1/2 \log_2 1/2 = 1$$

$$\begin{aligned} H(I, \text{Eye-colour}) &= \frac{3}{8} (-1 \log_2 1 - 0 \log_2 0) + \\ &\quad \frac{2}{8} (-0 \log_2 0 - 1 \log_2 1) + \\ &\quad \frac{3}{8} (-1/3 \log_2 1/3 - 2/3 \log_2 2/3) = 0.344 \end{aligned}$$

$$\begin{aligned} H(I, \text{Hair-colour}) &= \frac{2}{8} (-1/2 \log_2 1/2 - 1/2 \log_2 1/2) + \\ &\quad \frac{6}{8} (-3/6 \log_2 3/6 - 3/6 \log_2 3/6) = 1 \end{aligned}$$

$$\begin{aligned} H(I, \text{Height}) &= \frac{2}{8} (-1 \log_2 1 - 0 \log_2 0) + \\ &\quad \frac{4}{8} (-1/2 \log_2 1/2 - 1/2 \log_2 1/2) + \\ &\quad \frac{2}{8} (-0 \log_2 0 - 1 \log_2 1) = 0.5 \end{aligned}$$

ID3: example (11)

$$G(I, \text{Eye-colour}) = 1 - 0.366 = \mathbf{0.634}$$

$$G(I, \text{Hair-colour}) = 1 - 1 = 0$$

$$G(I, \text{Height}) = 1 - 0.5 = 0.5$$



	Hair-colour	Height	Class
E4	Brown	Medium	C-
E5	Brown	Tall	C+
E7	Blonde	Low	C-

ID3: example (12)

$$H(\text{Green}) = \underbrace{-1/3 \log_2 1/3}_{(E_5)} - \underbrace{2/3 \log_2 2/3}_{(E_4, E_7)} = 0,918$$

$$H(\text{Green}, \text{Hair-colour}) = \mathbf{1/3} (-\log_2 1 - 1 \log_2 1) + \mathbf{2/3} (-1/2 \log_2 1/2 - 1/2 \log_2 1/2) = 2/3$$

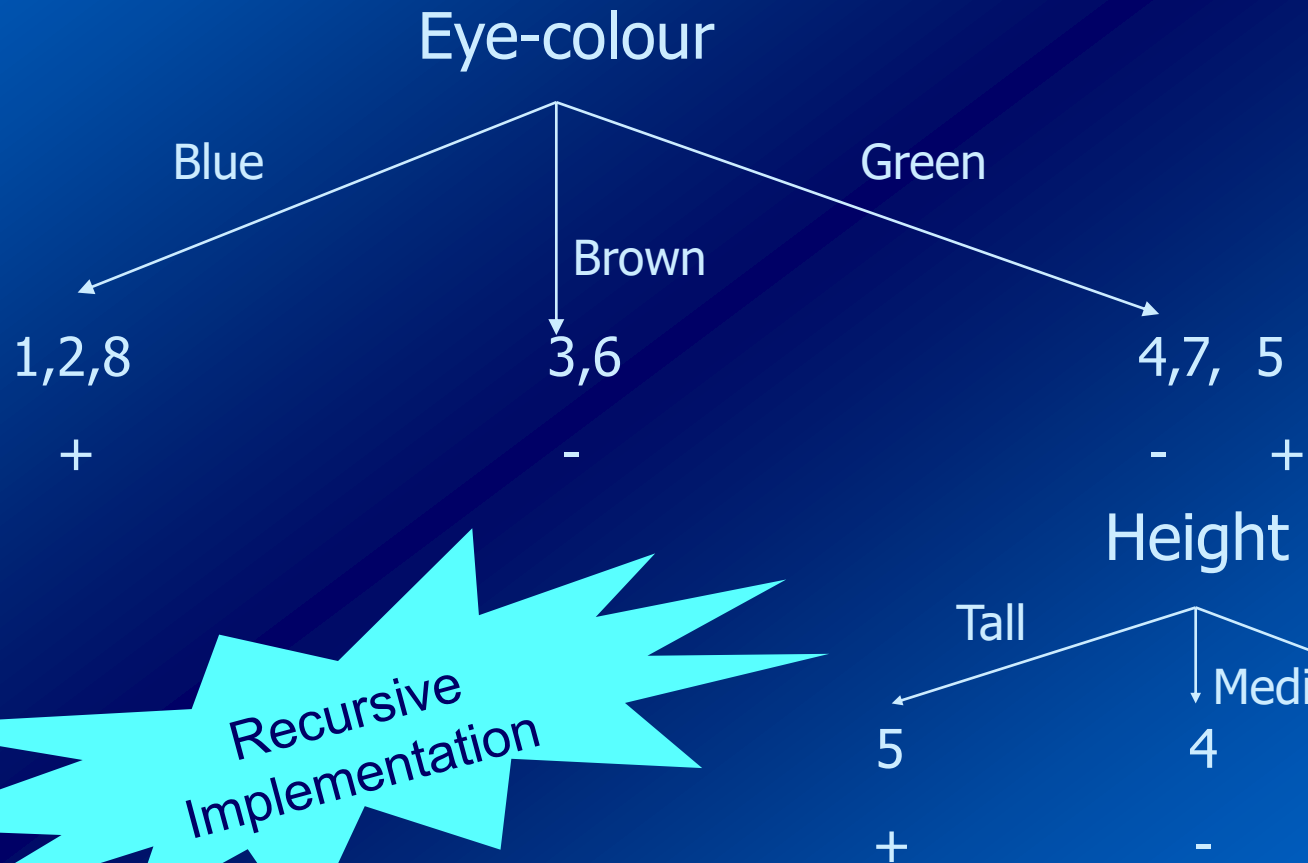
Blonde
↖

$$H(\text{Green}, \text{Height}) = 1/3 (-0 \log_2 0 - 1 \log_2 1) + 1/3 (-1 \log_2 1 - 0 \log_2 0) + 1/3 (-0 \log_2 0 - 1 \log_2 1) = 0$$

$$G(\text{Green}, \text{Hair-colour}) = 0,918 - 0,666 = 0,252$$

$$G(\text{Green}, \text{Height}) = 0,918 - 0 = \mathbf{0,918}$$

ID3: example



Recursive
Implementation

ID3: example (14)

Eye-colour = Blue \rightarrow C+

Eye-colour = Brown \rightarrow C-

Eye-colour = Green \wedge Height = Tall \rightarrow C+

Eye-colour = Green \wedge Height = Medium \rightarrow C-

Eye-colour = Green \wedge Height = Low \rightarrow C-



Transcription in a set
of rules

ID3: example (15)

Optimizing the set of rules

$(\text{Eye-colour} = \text{Blue}) \text{ or } (\text{Eye-colour} = \text{Green} \wedge \text{Height} \neq \text{Tall}) \rightarrow C+$

$(\text{Eye-colour} = \text{Brown}) \text{ or } (\text{Eye-colour} = \text{Green} \wedge \text{Height} = \text{Tall}) \rightarrow C-$



Leo Breiman,
American, 1928-2005

The CART solution




Jerome Friedman,
American

Developed 1974-1984 by 4 statistics professors: Leo Breiman (Berkeley), Jerry Friedman (Stanford), Charles Stone (Berkeley), Richard Olshen (Stanford). Distributed by Salford Systems <http://salford-systems.com/>

- Just perform Binary trees
- Unifies the categorical and continuous response under the same framework.
 - Classification tree
 - Regression tree
- Any kind of explanatory variable
- Split criterion: Impurity of the node
- Post pruning (without stop criterion)
- Delivers honest estimates of the quality of a tree

Selection of the optimal partition

Maximize the decrement of impurity between the parent and its children



$$\Delta i(t) = i(t) - \frac{n_{tl}}{n_t} i(t_l) - \frac{n_{tr}}{n_t} i(t_r)$$

Impurity of a node

For categorical responses:

- Gini

$$i(t) = \sum_{j \neq i} p(j/t) p(i/t) = 1 - \sum_j p_j^2$$

- Information (Entropy)

$$i(t) = - \sum_j p(j/t) \log_2 p(j/t)$$

For continuous responses:

- Variance

$$i(t) = \frac{\sum_{i \in t} (y_i - \bar{y}_t)^2}{n_t}$$

Missclassification tax

Gini(Node labelled as C) = $p(c|I) * (1-p(c|I))$ two classes

Gini(Node labelled as C) = $1 - \sum_{c' \in P, c' \neq C} p(c'|I)^2$, more than 2 classes

Selection of the optimal tree (from the maximum tree)

We need to define how good (or bad) is a tree:

Obvious criterion = *Probability of misclassification (= cost of the tree)*

How can we compute the Cost of the tree:

We assign every leave to the response class with maximum $p(j/t)$

$$p(j/t) = \frac{p(j,t)}{p(t)} = \frac{\pi_j \frac{n_{jt}}{n_j}}{\frac{n_{jt}}{n}}$$

If classes j autorepresentatives

$$\text{if } \pi_j = \frac{n_j}{n} \quad p(j/t) = \frac{n_{jt}}{n_j}$$

Cost of the tree

of a node: $r(t) = 1 - \max_j p(j / t)$

If there are missclassification costs

$$r(t) = \min_i \sum_j c(i / j) p(j / t)$$

Cost of the tree:

(decreasing with size)

$$R(T) = \frac{\sum_{t \in \tilde{T}} p(t) r(t)}{r(\text{root})} \times 100$$

Criterion to optimize: $\text{Min } R(T)$

Cost of a regression tree

- In a regression tree, every individuals is assigned with the mean value of the leave.
- Cost of the tree = residual variance

$$r(t) = \frac{1}{n_t} \sum_{i=1}^{n_t} (y_{it} - \bar{y}_t)^2$$

Penalization for complexity

Criterion to optimize:

$$\text{Min} \left(R(T) + \alpha |T| \right)$$

Where α is the *complexity parameter*, expresses the penalization for building large trees

Run with increasing complexity and

Get a sequence of optimal trees of increasing size....

Evaluate and select the best

Unbalanced Data

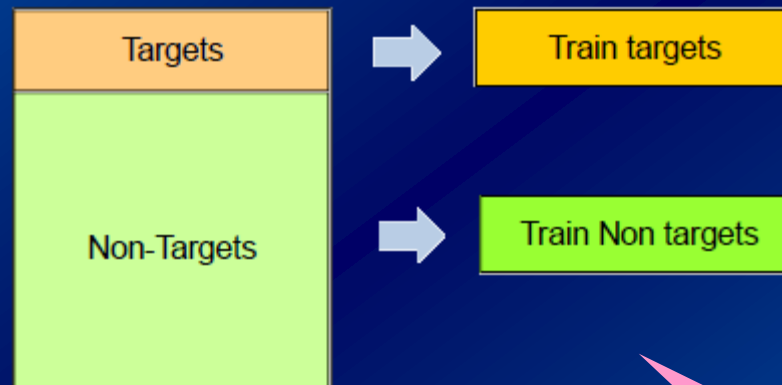
- Unequal class frequencies
 - 99% don't buy, 1% buy
 - 90% healthy, 10% disease
- Getting a majority class of 97% accuracy is useless
- Decision Trees assume balanced classes

(minority class hardly become a majority in a node)

- Work with balanced training datasets

Unbalanced training Data

Bad performance



Otherwise, weight non targets:

$$\sum_{\text{targets}} \text{weights} = \sum_{\text{non targets}} \text{weights}$$

Keep
Proportions

After Learning: Using

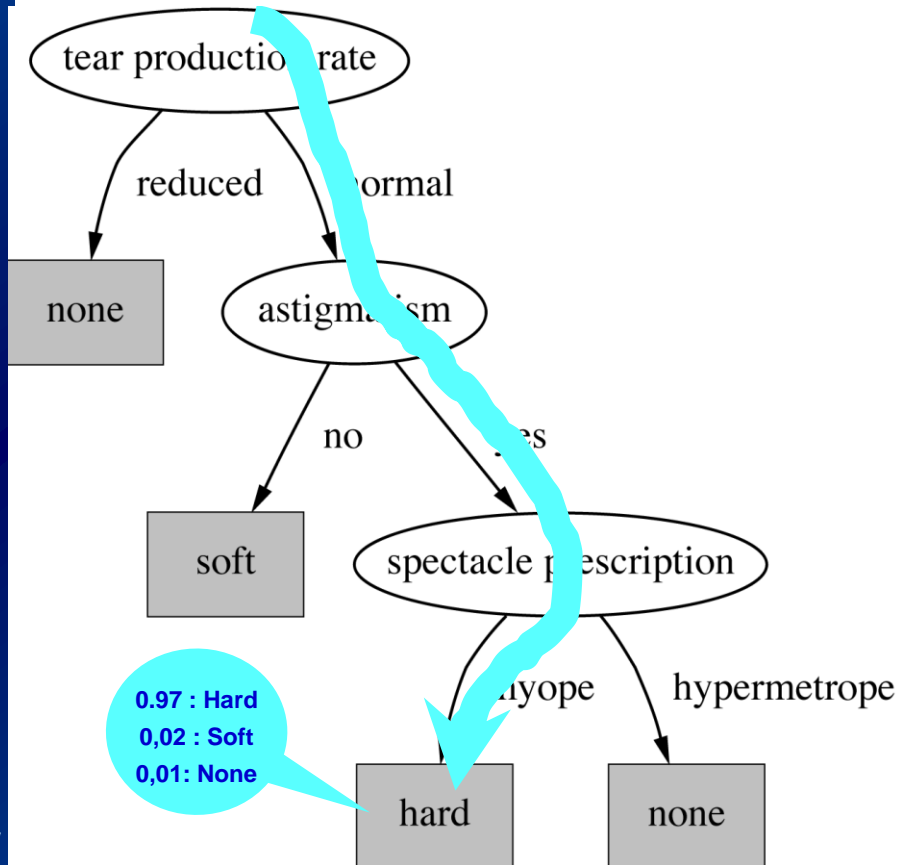
- Use the inducted DT to assign the class of a new object:

PREDICTION

- Evaluate node-queries over the object and descend to a leaf
- Assign the label of the leaf as the class of the object

Eventually, consider the purity of the leaf as the probability of the assignment

John: (normal tear, Astigmatism, myope)



OUTPUT: John requires hard lents

With a probability of 0.97

Validation in Decision Trees

Missclassification rate

Confusion matrix

Cross-validation

ROC curves: (DT provides single points for ROC space)

use probability of assignment to run ROC

or move some parameter (cp, minleaf....)

Confusion matrix and error rate

	Predicted class YES	Predicted class NO
Real class YES	n_{TP}	n_{FN}
Real class NO	n_{FP}	n_{TN}

Type-I
error

Type-II
error

$$\text{Error rate} = \frac{n_{FN} + n_{FP}}{n}$$

Missclassification rate depends on DT parameters

cpTable

	CP	nsplit	rel error	xerror	xstd
1	0.30312557	0	1.0000000	1.0405435	0.02080410
2	0.04100900	1	0.6968744	0.7092905	0.01917514
3	0.02817402	2	0.6558654	0.6547873	0.01854656
4	0.02200789	5	0.5596869	0.6307159	0.01895873
5	0.01433087	6	0.5376790	0.6151649	0.01925346
6	0.01419604	7	0.5233481	0.6048818	0.01967947
7	0.01372608	8	0.5091521	0.6034818	0.01967728
8	0.01133242	9	0.4954260	0.5893790	0.01932271
9	0.01000000	11	0.4727612	0.5694968	0.01876309

Cp: complexity parameter

Rel error: relative missclassification rate of the tree at each cp

Xerror: Crossvalidated missclassification error

Xstd: standard deviation of crossvalidated missclassification error

Pros and cons of DT

- Easy to interpret
- Branches provide assignment rules (Results operationalized)
- Branches provide good support to decision making
- Automatic detection of complex interactions among variables
- Computationally efficient
- But.... Rough approach to the response (approached by leaps (discontinuities))
- But... worse performance than other classifiers

rpart *package:rpart*

Recursive Partitioning and Regression Trees

```
rpart(formula, data, weights, na.action = na.rpart, method, parms, control, cost, ...)
```

formula: a formula, as in the 'lm' function.

data: an optional data frame in which to interpret the variables named in the formula

weights: optional case weights.

na.action: The default action deletes all observations for which 'y' is missing, but keeps those in which one or more predictors are missing.

method: one of "anova", "poisson", "class" or "exp". If 'method' is missing then the routine tries to make an intelligent guess.

```
parms=list(prior=c(.85,.15), split='information')
```

```
control=rpart.control(minsplit=20, minbucket=round(minsplit/3), cp=0.01,maxcompete=4,  
  maxsurrogate=5, usesurrogate=2, xval=10, maxdepth=30, ...)
```

minsplit: the minimum number of observations that must exist in a node, in order for a split to be attempted.

minbucket: the minimum number of observations in any terminal '<leaf>' node.

cp: complexity parameter.

maxcompete: the number of competitor splits retained in the output.

maxsurrogate: the number of surrogate splits retained in the output.

usesurrogate: 1= use surrogates, in order, to split subjects missing the primary variable; if all surrogates are missing the observation is not split. 2= if all surrogates are missing, then send the observation in the majority direction.

xval: number of cross-validations

maxdepth: Set the maximum depth of any node of the final tree, with the root node counted as depth 0